

Creo Advisors — AI Value-Creation Marketing Plan

“AI value is a map, not a multiplier.”

Draft marketing plan prepared by the ENT 105 team for Creo Advisors

2 June 2026

The big idea. Private equity’s hottest play is the “AI roll-up” — buy fragmented, labor-heavy services businesses and use AI to cut cost and expand margin. The hype says AI replaces labor wholesale (“1000:1”). The verified reality: most PE firms can’t show an AI return at all. The truth in between is that **the AI handoff ratio is a *map across functions*, not a single multiplier** — a few coordination functions collapse at $\sim 10,000:1$, the knowledge-work middle amplifies at $\sim 3:1$, and frontline judgment work stays $\sim 1:1$. **Creo’s market is telling PE buyers which is which** — so they don’t overpay for the average, or gut the business chasing it. The lead product is the **AI-Automatability Diagnostic**: an independent, buyer-side read that generates qualified leads and reprices deals.

Contents

1	Objective	2
2	The opportunity & why now	2
3	The core insight: the handoff ratio is a map	2
4	Competitive landscape & the gap	4
5	Positioning & value proposition	4
6	Target personas	5
7	The flagship offer: the AI-Automatability Diagnostic	5
8	Campaign & go-to-market	6
9	Research plan (primary & secondary)	6
10	Budget (illustrative, annual — boutique scale)	7
11	KPIs	7

12 12-month phasing & course milestones	7
13 Evidence base (cited)	7

1 Objective

Position Creo to win AI-value-creation mandates from PE sponsors and their portfolio companies, and to **generate qualified leads** for that work — toward the 2027 revenue goal to be confirmed at the Discovery meeting. The mechanism: a free, rigorous *AI-Automatability Diagnostic* that draws PE buyers in, qualifies them, and hands warm leads to the partners.

2 The opportunity & why now

- **The AI roll-up is a real, funded 2024–26 wave.** Articulated by Elad Gil and scaled by General Catalyst (Eudia/legal, Titan/IT, Long Lake — which took Amex GBT private for \$6.3B), Thrive Holdings (\$1B+, OpenAI-backed), Khosla, Lightspeed, and others, across legal, accounting, IT/MSP, BPO/call centers, healthcare, and property management.
- **And the model just went private-sector, on the record.** In June 2026, ex-DOGE staffers Nate Cavanaugh and Justin Fox launched **Special** — an a16z-led firm (backers include Steve Davis, Valor’s Antonio Gracias, ex-xAI CFO Anthony Armstrong, Coinbase’s Brian Armstrong, and Palantir’s Shyam Sankar) whose thesis is, in its own words, “buy a company, automate the overhead, keep the margin.” Its first target is a *nurse-heavy eldercare business in Texas* (1,400+ patients, hundreds of nurses). By its own framing it is *a thesis with funding — not yet a track record*: the exact moment a buyer-side, rigor-over-hype read is worth the most.
- **But the results aren’t there — which is the opening.** BCG (2026): *most PE firms cannot show meaningful AI returns*; only **5% of companies capture substantial value while 60% lag**. An MIT-linked study found **~95% of enterprise GenAI pilots fail**; Klarna publicly reversed its AI-only support push and began *re-hiring humans*. The 2025 data keeps converging: S&P Global / 451 Research found enterprises abandoning the *majority* of their AI initiatives leapt **17%→42%** in a year ($n=1,006$), and IBM’s survey of **2,000 CEOs** found only **25%** of AI initiatives delivered the expected ROI and only **16%** scaled enterprise-wide. And at national scale, the very playbook the *Special* founders came from is contested: DOGE claims **~\$160B** saved, while one analysis (the nonpartisan Partnership for Public Service) estimates its actions will *cost* taxpayers **~\$135B** this fiscal year once paid leave, re-hiring, and lost productivity are counted — the cut that books as savings and returns as cost.
- **So buyers are flying blind into the hottest strategy in PE.** They are pricing deals on hype ratios with no rigorous way to tell what’s real. That gap — between the pitch and the durable truth — is the market Creo can own.

3 The core insight: the handoff ratio is a map

The single idea the whole offering rests on, and the reason a rigorous diagnostic is worth paying for:

Function type	Handoff ratio	What it is
Collapse	~10,000:1	Pure coordination / matching / information-routing — the “dispatcher.” Near-total automation. <i>Precedent: Uber/DoorDash replaced the human dispatch layer of an entire industry with one matching algorithm.</i>
Amplify	~3:1	Knowledge work where AI makes a skilled operator several times more effective (human still in the loop).
Level-up	—	Routine work where AI raises the floor / lifts the average performer.
Keep human	~1:1	Judgment, physical, relationship, trust, and regulated work — not durably automatable (Uber still needs the drivers).

Why this is the wedge. The hype prices every function at ~1000:1 and overpays. The skeptic prices everything at ~3:1 and misses the rare 10,000:1 prize. **Both are wrong, because a business is a *portfolio* of functions with wildly different ratios.** The value — and the risk — live in the mix, and almost no one measures it. Creo does.

(Ratios are illustrative orders of magnitude; the real figures are function- and industry-specific — which is precisely what the Diagnostic measures.)

The *shape* of the map is evidenced, not asserted. Peer-reviewed causal studies show AI’s effect is heterogeneous and task-dependent — not a uniform multiplier. In a ~5,000-agent field study (Brynjolfsson, Li & Raymond, *Quarterly Journal of Economics* 2025), generative AI lifted support productivity **+34% for novices but ~0 for the most experienced** — it raises the floor; it doesn’t replace the expert. In a 758-consultant randomized trial (Dell’Acqua et al., *Organization Science* 2025), AI *inside* its “jagged frontier” raised quality **40%+** — yet on a task just *outside* it, AI users were **19 points less likely** to be correct. And a 2025 METR trial found experienced developers were **19% slower** with AI while believing they were faster — a caution on self-reported gains. Same tools, opposite results, by function. That is the map.

The second map: who can be re-selected. The function map says *what* could be automated; it doesn’t say whether *this* business can pull it off. The durable value comes from a workforce **swap, not a purge**: keep and amplify the operators who truly understand the business, collapse the pure-coordination layer, and trade non-adopters for AI-fluent talent — without breaking quality. How large a swap a business can sustain is itself a ratio, specific to the company and its management. Microsoft is the public template: it cut 15,000+ in 2025 while *mandating* AI use and shielding its AI teams — reselection toward AI-fluent talent, not 1000:1 replacement. **The tell:** when a company sheds thousands and rehires a sizable *fraction* as AI adopters, that fraction is a readout of the real automation ratio — amplification (low single digits), not replacement (1000:1).

The market is starting to confirm the swap. Gartner forecasts that **by 2027, 50% of companies that cut staff citing AI will re-hire** for similar service roles under new titles, judging today’s AI “not mature enough” to replace human judgment and empathy (a near-term call, not a permanent one). And the amplify-not-gut model is already being bought: in acquiring CX outsourcer PartnerHero (2024), AI firm Crescendo *absorbed* ~3,000 “human-in-the-loop” agents rather than cutting

them — AI to free people for higher-value work (the deal parties’ own framing).

4 Competitive landscape & the gap

The field is crowded — but it sorts into three seats, and almost every player sells the AI *build* its diligence funnels into. (Offerings below are firms’ own positioning, attributed not asserted.)

Player	Named AI-for-PE offering	Seat	Sells the build?
Bain	AI folded into a 5-area “Full-Potential” diligence (no standalone read)	Diligence (embedded)	Yes
BCG	“Deploy / reshape / invent”; 100-day diagnostic, portfolio playbooks	Build / platform	Yes
EY-Parthenon	Named “AI due diligence” (market / product / R&D / operations)	Diligence	Yes (EY.ai)
RSM	AI-readiness: 6 prerequisites, “EBITDA in 6 months” filter	Diligence → build	Yes (\$1B AI)
AlixPartners	“Outside-In Due Diligence” + expert-based AI DD (<i>nearest analogue</i>)	Diligence	Yes (firm-level)
A&M	“A&M Assist” / “DiligenceGPT” — proprietary GenAI platform	Platform / build	Yes
FTI	“Three plays”: buy-side + sell-side narrative + build	Diligence + build	Yes
In-house (Vista, EQT)	“Agentic AI Factory” / “Mother-brain” — captive portfolio platforms	Platform (internal)	Yes (captive)

The gap — the seat no one owns. Every credible incumbent either **sells the AI build its diligence leads into** (Bain, BCG, EY, RSM, A&M, FTI; AlixPartners at firm level) or runs a **captive portfolio platform** not sold to buyers (Vista, EQT); talent firms (e.g. Korn Ferry) only *observe* the trend. The closest structural analogues to Creo — AlixPartners’ Outside-In Due Diligence and EY-Parthenon’s four-dimension AI diligence — are genuinely buyer-side reads, but both sit *inside firms that also build AI for portfolio companies*, so firm-level independence doesn’t hold. The **independent, buyer-side, diligence-grade, sells-no-AI-systems** seat — rigor over hype, no build to upsell — is **open**. That is the seat Creo takes.

5 Positioning & value proposition

Positioning. *Creo is the independent, buyer-side authority on what AI can actually do in a deal — mapping the real automation ratio function by function, so sponsors capture the durable value and avoid the deals that blow up.*

Three reasons to believe (all defensible, none a novelty claim):

- **Independent — nothing to sell on the other side.** Creo sells no AI systems *and no build*. The verified pattern across the field (strategy firms, Big Four, ops shops): AI “diligence” is a gateway to the build they monetize. Creo’s read on *what to automate* funnels into nothing it sells — in a market full of AI claims, that neutrality is the scarce asset.

- **Rigorous & transparent.** Grounded in the established task-automation research (Frey–Osborne’s automation-resistant task classes; the “exposure $\neq$ automation” caveat from the OpenAI/Science study) — every call reconciles and is shown, not asserted.
- **Amplify, don’t gut.** The durable play is a *swap, not a purge* — keep and amplify the operators who run the business, collapse only the busywork, and re-select for AI-fluent talent; the gut-and-replace failure mode (deploy-and-fire) is exactly why most AI programs show no return.

6 Target personas

Persona	Pain / trigger	Message
Deal Partner / MD (<i>pre-deal</i>)	Underwriting an AI-rollup thesis; can’t tell real from hype; IC scrutiny	“Price on the real ratio, not the pitch — before you overpay.”
Operating Partner (<i>champion</i>)	Must show AI returns to LPs; portfolio pilots stalling	“Most AI programs show no return. We map where the durable value actually is.”
Portfolio CEO (<i>operator</i>)	Pressure to “do AI”; fear of breaking the business	“Amplify your best people and collapse the busywork — without gutting what works.”
LP / IR (<i>amplifier</i>)	Differentiating the fund’s value-creation story	“Back your AI story with diligence you can defend.”

7 The flagship offer: the AI-Automatability Diagnostic

A fast, independent assessment along **two axes**: (1) a *function* map — which of a target’s functions fall into Collapse / Amplify / Level-up / Keep-human; and (2) a *reselection* read — whether the business has operators worth amplifying and management able to execute the swap (cull non-adopters, bring in AI-fluent talent) without breaking quality. Together they roll into a **blended, durable value estimate** and **what the deal is actually worth** — versus the seller’s headline.

Outputs: the function-by-function ratio map; a **reselection-readiness** score (amplifiable talent + management’s capacity to run the swap); the “real vs. hyped value” gap; the amplify/collapse/keep priorities; and the deals/functions to walk away from.

Offer & pricing ladder (free proof $\rightarrow$ recurring relationship):

1. **Free:** a gated “AI Reality Report” and a complimentary mini-Diagnostic on one target — the lead magnet and the proof.
2. **Productized:** a fixed-fee full Diagnostic on a deal or a portfolio company.
3. **Engagement:** the value-creation/amplification roadmap built on the Diagnostic (Creo orchestrates; vendors build).
4. **Recurring:** the trusted AI read across the sponsor’s portfolio.

Closing the loop — measured proof. Post-deal, Creo verifies that the projected durable value was actually realized, using its value-creation measurement methodology (the Value-Creation Bridge) to separate real operational gains from value merely transferred. The Diagnostic’s promise is held to account: diligence $\rightarrow$ amplification $\rightarrow$ *measured* result.

8 Campaign & go-to-market

A boutique-scale *report* → *amplify* → *engage* → *convert* engine:

Stage	Asset	Mechanic
Anchor	The AI Reality Report (annual)	Creo’s proprietary take on the real ratios + failure rates; gated ⇒ leads
Amplify	Founder-led LinkedIn + short pieces	Contrarian, evidence-based hooks (below); inbound/SEO
Engage	“Inside the Reality Report” roundtable/webinar	Converts readers to conversations with target sponsors
Convert	Complimentary mini-Diagnostic	The qualified-lead handoff to the partners

Signature content hooks (contrarian = attention + credibility):

- “*AI value is a map, not a multiplier*” — the function-ratio thesis.
- “*The 10,000:1 function hiding in your target*” — find the coordination layer.
- “*The AI layoffs aren’t what they look like*” — Microsoft cut 15,000+ in 2025 while *mandating* AI use (tied to performance reviews) and shielding its own AI teams from the freeze: reselection toward AI-fluent talent, not 1000:1 replacement.
- “*The rehire is the tell*” — when a giant lays off thousands and quietly rehires a sizable share as AI adopters, that fraction *is* the real automation ratio.
- “*95% of AI pilots fail. Here’s the 5%.*”
- “*a16z just funded DOGE-for-PE*” — the *Special* launch as the live test of the thesis, and the \$135B question it can’t yet answer.
- “*The model bought a company full of nurses*” — when “automate the overhead” meets the most *keep-human* function there is, *which* overhead is the whole game.

Channels & tools: LinkedIn (primary, founder-led); email nurture and CRM in **HubSpot** (inbound stack); SEO around “AI value creation / AI due diligence”; small sponsor-targeted events; a fast, crawlable site with a real lead-capture CTA. Selective **outbound** via warm sponsor introductions and PE-forum speaking.

9 Research plan (primary & secondary)

Secondary (done / ongoing): the AI-rollup deal landscape, the competitive field, the automation-research lineage, and the failure-rate evidence (see Evidence base). **Primary** (to run): short interviews/surveys with 5–10 PE operating and deal partners on (a) their AI-rollup intentions, (b) where their pilots stall, and (c) what they’d pay to de-risk — to validate the personas, the pain, and the willingness-to-pay before the final plan.

10 Budget (illustrative, annual — boutique scale)

Line item	Illustrative \$
Annual AI Reality Report (research, analysis, design, PR)	25,000–40,000
Website (crawlable, lead capture, case examples)	10,000–20,000
HubSpot (marketing + CRM) + nurture	6,000–12,000 / yr
LinkedIn / content production	12,000–24,000 / yr
Roundtables / events (2–3 / yr)	10,000–20,000
Total cash (ex. partner time)	~\$65,000–115,000 / yr

Most leverage is partner *time* (the report’s point of view, LinkedIn, roundtables), not cash — the right shape for a boutique. Figures illustrative.

11 KPIs

Awareness: report downloads; LinkedIn growth/engagement; organic traffic. **Interest:** MQLs; webinar/roundtable attendance. **Intent:** mini-Diagnostic requests (*key leading indicator*). **Conversion:** qualified opportunities; diagnostics → paid engagements; revenue. **Efficiency:** CAC and LTV (and LTV:CAC) — the boutique’s North-Star unit economics.

12 12-month phasing & course milestones

Quarter	Focus	Course tie-in
Q1	Foundations: site + capture, HubSpot, personas/messaging, run the report survey	Discovery / Wk-4 presentation
Q2	Launch the AI Reality Report + founder LinkedIn cadence	Mid-term draft
Q3	Roundtable + outbound + complimentary Diagnostics	HBS case lessons applied
Q4	Convert diagnostics to engagements; report CAC/LTV; scope Year 2	Wk-13 final (mgmt)

13 Evidence base (cited)

- AI-rollup thesis (Gil) — TechCrunch, Jun 2025. / GC & Long Lake–Amex GBT \$6.3B — PitchBook & BusinessWire, 2026. / Thrive Holdings / OpenAI stake — Newcomer; Bloomberg.
- Most PE can’t show AI returns; 5% / 60%; deploy vs. reshape — BCG, “Inside the AI-First Private Equity Firm,” 2026; “The Widening AI Value Gap.”
- ~95% of GenAI pilots fail — MIT (via Fortune, Aug 2025). / Klarna re-hiring — Entrepreneur, 2025.
- Microsoft 2025: >15,000 roles cut (May ~6,000 + July ~9,000) — CNBC, Jul 2025; plus a ~7% US voluntary-buyout offer — HCAmag, 2026. AI use made “no longer optional” and tied to performance reviews (Liuson memo) — Business Today / HR Grapevine, Jun 2025. AI/Copilot teams shielded from the 2026 hiring freeze — CNBC, Apr 2026.
- Automation-resistant task classes; 47% high-risk (not a forecast) — Frey & Osborne (2013). / “exposure ≠ automation”; ~80% of workers, ≥10% of tasks LLM-exposed — OpenAI “GPTs are GPTs” (Science, 2024).
- Incumbents — Bain (industry/PE pages, ~80% of largest PE); EY (PE value-creation); Vista (Rule-of-60, 85+ portcos via Bain PE Report).

- *Special* (private-sector AI roll-up; a16z-led; founders Nate Cavanaugh & Justin Fox; first target a Texas eldercare business, 1,400+ patients) — Bloomberg & The Next Web, Jun 2026; a16z announcement. / DOGE ~\$160B claimed savings vs. ~\$135B estimated cost this FY (paid leave, re-hiring, lost productivity; direct layoff savings only ~\$38B over 10yr; excludes litigation and lost tax collection) — Partnership for Public Service, via CBS News & Fortune, Jul 2025.
- Amplify $\neq$ replace (peer-reviewed) — Brynjolfsson, Li & Raymond, “Generative AI at Work,” *QJE* 2025 (NBER w31161): +14% avg, +34% novices, ~0 experts. / Dell’Acqua et al., “Navigating the Jagged Technological Frontier,” *Organization Science* 2025 (HBS WP 24-013): +40% quality inside the frontier, -19 pts outside. / METR (Jul 2025): experienced developers 19% *slower* with AI (perception $\neq$ measured reality).
- Returns & adoption — S&P Global Market Intelligence / 451 Research, “Voice of the Enterprise: AI & ML 2025” (abandonment 17%→42%; $n=1,006$; self-reported). / IBM Institute for Business Value “CEO Study” (May 2025; 2,000 CEOs; 25% expected ROI, 16% scaled; vendor-published). / Federal Reserve FEDS Note (Apr 2026): ~18% of U.S. firms had adopted AI at year-end 2025.
- Re-hiring & the augment model — Gartner (Feb 2026): forecasts 50% of AI-attributed staff cuts re-hired by 2027 (analyst prediction). / Crescendo-PartnerHero (Oct 2024): AI firm absorbed ~3,000 human-in-the-loop agents (deal-party framing).
- Competitive field (firms’ own positioning, 2023–26) — Bain (“Full-Potential” diligence); BCG (“deploy / reshape / invent”); EY-Parthenon (named “AI due diligence,” 4 dimensions); RSM (AI-readiness, “EBITDA in 6 months”); AlixPartners (“Outside-In Due Diligence”); A&M (“A&M Assist” / “DiligenceGPT”); FTI (“three plays”); Korn Ferry (“AI Operating Partner”); in-house: Vista (“Agentic AI Factory”), EQT (“Motherbrain”).

Honest caveats (this is a rigor-over-hype brand — it must hold itself to it). The illustrative ratios (10,000:1, 3:1, 1000:1) are orders-of-magnitude, not constants — the real figures are what the Diagnostic measures. Vendor/survey market figures (Bain’s “80%,” EY allocations) are self-reported — attribute, don’t assert. The “Uber/DoorDash” point is used as factual *precedent* (dispatch-layer automation), not a borrowed thesis. The Microsoft case is now verified and sourced (the 2025 cuts and the “AI is no longer optional” performance-review memo); the content rests on those documented facts and frames it as reselection-for-amplification, without overstating cause. The “white space” is a defensible read, not a proven vacancy — the incumbents are strong; Creo wins on focus, independence, and credibility, not scale.

Draft marketing plan prepared by the ENT 105 team for Creo Advisors. Figures and timelines illustrative, to be finalized with management. All analyses are independent and transparent — credibility is shown, not claimed.